

# Take Home Assessment - Part 2 - Segmentation Model - Gauravi Patankar (gsp2137@columbia.edu)

## Project Overview & Business Context

- The dataset comprises population survey data containing rich information on individuals' **demographic, educational, employment, and socioeconomic characteristics**.
  - The objective is to perform **unsupervised socio-demographic segmentation** and generate **target personas** that represent unique population groups.
  - Each segment will capture a coherent socio-demographic profile defined by attributes such as **age, education, occupation type, income proxies**, and **household structure**, among others.
- 

## Assumption

Since the details of the marketing campaign are **not specified**, this segmentation is treated as **exploratory** in nature, aimed at discovering **general**, meaningful, interpretable clusters rather than optimizing for a predefined campaign outcome.

---

## Business Relevance

A segmentation model is only as valuable as the clarity of the business questions it answers.

While the dataset provides comprehensive socio-demographic information, the **absence of campaign-specific context** limits the ability to align clusters with direct marketing actions.

To make the segmentation more targeted and actionable, the following **strategic questions** would ideally be clarified with the marketing or business team:

1. **Target Audience** – Which segment of the population is the campaign intended for? (e.g., working adults, families, students)
2. **Granularity** – Do we need a few broad, easily interpretable personas, or more granular segments for precision targeting?
3. **Ethical Boundaries** – Are there sensitive variables (e.g., race, nationality) that should be excluded from modeling for compliance or ethical reasons?

4. **Actionability** – How will marketing leverage these clusters (for message tailoring, media planning, or outreach prioritization)?

Addressing these questions would make the business objective explicit, helping focus the analysis on the most decision-relevant variables.

## Importing Libraries

```
In [1]: # Core libraries
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing, clustering, metrics, decomposition
from sklearn.preprocessing import RobustScaler, OrdinalEncoder, OneHotEncoder
from scipy.stats import chi2_contingency
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering, MiniBatchKMeans
from sklearn.mixture import GaussianMixture
from sklearn.metrics import (
    silhouette_score,
    calinski_harabasz_score,
    davies_bouldin_score
)

# Dimensionality reduction for visualization (optional)
from sklearn.manifold import TSNE

# Optional: for winsorization and outlier capping
from scipy.stats.mstats import winsorize

# Optional advanced visualization
import plotly.express as px

# General setup
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: # Random seed for reproducibility
np.random.seed(42)

# Display options
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 100)
pd.set_option('display.float_format', '{:.3f}'.format)
```

## Data

```
In [3]: # Reading column names
with open("census-bureau.columns") as f:
    columns = [line.strip() for line in f.readlines()]

# Reading data
df = pd.read_csv("census-bureau.data", names=columns, header=None)
```

```
In [4]: # Viewing data
print(df.shape)
df.head()
```

(199523, 42)

Out[4]:

	age	class of worker	detailed industry recode	detailed occupation recode	education	wage per hour	enroll in edu inst last wk	marital stat	ir
0	73	Not in universe	0	0	High school graduate	0	Not in universe	Widowed	univ (
1	58	Self- employed- not incorporated	4	34	Some college but no degree	0	Not in universe	Divorced	Cons
2	18	Not in universe	0	0	10th grade	0	High school	Never married	univ (
3	9	Not in universe	0	0	Children	0	Not in universe	Never married	univ (
4	10	Not in universe	0	0	Children	0	Not in universe	Never married	univ (

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 199523 entries, 0 to 199522
```

```
Data columns (total 42 columns):
```

#	Column	Non-Null Count		Dtype
----	-----	-----	-----	-----
0	age	199523	non-null	int64
1	class of worker	199523	non-null	object
2	detailed industry recode	199523	non-null	int64
3	detailed occupation recode	199523	non-null	int64
4	education	199523	non-null	object
5	wage per hour	199523	non-null	int64
6	enroll in edu inst last wk	199523	non-null	object
7	marital stat	199523	non-null	object
8	major industry code	199523	non-null	object
9	major occupation code	199523	non-null	object
10	race	199523	non-null	object
11	hispanic origin	198649	non-null	object
12	sex	199523	non-null	object
13	member of a labor union	199523	non-null	object
14	reason for unemployment	199523	non-null	object
15	full or part time employment stat	199523	non-null	object
16	capital gains	199523	non-null	int64
17	capital losses	199523	non-null	int64
18	dividends from stocks	199523	non-null	int64
19	tax filer stat	199523	non-null	object
20	region of previous residence	199523	non-null	object
21	state of previous residence	199523	non-null	object
22	detailed household and family stat	199523	non-null	object
23	detailed household summary in household	199523	non-null	object
24	weight	199523	non-null	float64
25	migration code-change in msa	199523	non-null	object
26	migration code-change in reg	199523	non-null	object
27	migration code-move within reg	199523	non-null	object
28	live in this house 1 year ago	199523	non-null	object
29	migration prev res in sunbelt	199523	non-null	object
30	num persons worked for employer	199523	non-null	int64
31	family members under 18	199523	non-null	object
32	country of birth father	199523	non-null	object
33	country of birth mother	199523	non-null	object
34	country of birth self	199523	non-null	object
35	citizenship	199523	non-null	object
36	own business or self employed	199523	non-null	int64
37	fill inc questionnaire for veteran's admin	199523	non-null	object
38	veterans benefits	199523	non-null	int64
39	weeks worked in year	199523	non-null	int64
40	year	199523	non-null	int64
41	label	199523	non-null	object

```
dtypes: float64(1), int64(12), object(29)
```

```
memory usage: 63.9+ MB
```

## Data Cleaning

### 1. Duplicates

```
In [6]: df.duplicated().sum()
```

```
Out[6]: np.int64(3229)
```

```
In [7]: df.drop_duplicates()
```


Out [7]:

	age	class of worker	detailed industry recode	detailed occupation recode	education	wage per hour	enroll in edu inst last wk	marital stat
<b>0</b>	73	Not in universe	0	0	High school graduate	0	Not in universe	Widowed
<b>1</b>	58	Self-employed-not incorporated	4	34	Some college but no degree	0	Not in universe	Divorced
<b>2</b>	18	Not in universe	0	0	10th grade	0	High school	Never married
<b>3</b>	9	Not in universe	0	0	Children	0	Not in universe	Never married
<b>4</b>	10	Not in universe	0	0	Children	0	Not in universe	Never married
...	...	...	...	...	...	...	...	...
<b>199518</b>	87	Not in universe	0	0	7th and 8th grade	0	Not in universe	Married-civilian spouse present
<b>199519</b>	65	Self-employed-incorporated	37	2	11th grade	0	Not in universe	Married-civilian spouse present
<b>199520</b>	47	Not in universe	0	0	Some college but no degree	0	Not in universe	Married-civilian spouse present
<b>199521</b>	16	Not in universe	0	0	10th grade	0	High school	Never married
<b>199522</b>	32	Private	42	30	High school graduate	0	Not in universe	Never married

196294 rows × 42 columns

## 2. Fixing data types

```
In [8]: plt.figure(figsize=(6, 3))
sns.countplot(x='num persons worked for employer', data=df, color='skyblue')
plt.title('Distribution of Number of Persons Worked for Employer', fontsize=
plt.xlabel('Number of Persons Worked For Employer', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



```
In [9]: df['num persons worked for employer'] = df['num persons worked for employer']
```

```
In [10]: df['veterans benefits'].value_counts()
```

```
Out[10]: veterans benefits
2      150130
0      47409
1       1984
Name: count, dtype: int64
```

```
In [11]: df["own business or self employed"].value_counts()
```

```
Out[11]: own business or self employed
0      180672
2       16153
1        2698
Name: count, dtype: int64
```

```
In [12]: df['veterans benefits'] = df['veterans benefits'].astype('category')
df["own business or self employed"] = df['veterans benefits'].astype('category')
```

```
In [13]: obj_cols = df.select_dtypes(include='object').columns
df[obj_cols] = df[obj_cols].apply(lambda col: col.astype('category'))
```

### 3. Dropping Columns

Dropping irrelevant columns for segmentation

```
In [14]: # Drop detailed columns
drop_cols = [
    'detailed occupation recode',
    'detailed industry recode',
    'weight',
    'label',
    'year'
]
df.drop(columns=[c for c in drop_cols if c in df.columns], inplace=True)
```

### 4. Simplify Categories

```
In [15]: # education
df['education'] = df['education'].astype(str).str.strip().str.lower()
edu_group_map = {
    'less than 1st grade': 'Below High School',
    '1st 2nd 3rd or 4th grade': 'Below High School',
    '5th or 6th grade': 'Below High School',
    '7th and 8th grade': 'Below High School',
    '9th grade': 'Below High School',
    '10th grade': 'Below High School',
    '11th grade': 'Below High School',
    '12th grade no diploma': 'Below High School',
    'high school graduate': 'High School Graduate',
    'some college but no degree': 'Some College',
    'associates degree-occup /vocational': 'Associate Degree',
    'associates degree-academic program': 'Associate Degree',
    'bachelors degree(ba ab bs)': 'Bachelors Degree',
    'masters degree(ma ms meng med msw mba)': 'Graduate/Professional Degree',
    'prof school degree (md dds dvm llb jd)': 'Graduate/Professional Degree',
    'doctorate degree(phd edd)': 'Graduate/Professional Degree',
    'children': 'Children'
}
df['education'] = df['education'].map(edu_group_map)
df['education'] = df['education'].astype('category')

df['education'].value_counts(dropna=False)
```

```
Out[15]: education
High School Graduate      48407
Children                  47422
Below High School        36691
Some College              27820
Bachelors Degree         19865
Associate Degree          9721
Graduate/Professional Degree  9597
Name: count, dtype: int64
```

```
In [16]: # detailed household and family stat
```



```

df['detailed household and family stat'] = (df['detailed household and famil
    .astype(str).str.strip().str.lower()
)
household_group_map = {
    # Householder
    'householder': 'Householder',
    'nonfamily householder': 'Householder',

    # Spouse/Partner
    'spouse of householder': 'Spouse/Partner',
    'spouse of rp of unrelated subfamily': 'Spouse/Partner',

    # Child
    'child <18 never marr not in subfamily': 'Child',
    'child 18+ never marr not in a subfamily': 'Child',
    'child 18+ ever marr not in a subfamily': 'Child',
    'child <18 ever marr not in subfamily': 'Child',
    'child 18+ spouse of subfamily rp': 'Child',
    'child <18 spouse of subfamily rp': 'Child',
    'child 18+ never marr rp of subfamily': 'Child',
    'child <18 never marr rp of subfamily': 'Child',
    'child 18+ ever marr rp of subfamily': 'Child',
    'child <18 ever marr rp of subfamily': 'Child',
    'child under 18 of rp of unrel subfamily': 'Child',

    # Grandchild
    'grandchild <18 never marr not in subfamily': 'Grandchild',
    'grandchild <18 ever marr not in subfamily': 'Grandchild',
    'grandchild 18+ never marr not in subfamily': 'Grandchild',
    'grandchild 18+ ever marr not in subfamily': 'Grandchild',
    'grandchild 18+ spouse of subfamily rp': 'Grandchild',
    'grandchild <18 never marr rp of subfamily': 'Grandchild',
    'grandchild <18 ever marr rp of subfamily': 'Grandchild',
    'grandchild 18+ ever marr rp of subfamily': 'Grandchild',
    'grandchild 18+ never marr rp of subfamily': 'Grandchild',
    'grandchild <18 never marr child of subfamily rp': 'Grandchild',

    # Other Relative
    'other rel 18+ ever marr not in subfamily': 'Other Relative',
    'other rel 18+ never marr not in subfamily': 'Other Relative',
    'other rel <18 never marr not in subfamily': 'Other Relative',
    'other rel <18 ever marr not in subfamily': 'Other Relative',
    'other rel 18+ spouse of subfamily rp': 'Other Relative',
    'other rel <18 spouse of subfamily rp': 'Other Relative',
    'other rel 18+ ever marr rp of subfamily': 'Other Relative',
    'other rel 18+ never marr rp of subfamily': 'Other Relative',
    'other rel <18 ever marr rp of subfamily': 'Other Relative',
    'other rel <18 never married rp of subfamily': 'Other Relative',
    'other rel <18 never marr child of subfamily rp': 'Other Relative',
    'other rel 18+ spouse of subfamily rp': 'Other Relative',
    'other rel 18+ ever marr rp of subfamily': 'Other Relative',

    # Non-relative/Secondary
    'secondary individual': 'Non-relative/Secondary',
    'rp of unrelated subfamily': 'Non-relative/Secondary',

```

```

# Group quarters
'in group quarters': 'Group Quarters'
}

df['detailed household and family stat'] = (
    df['detailed household and family stat']
    .map(household_group_map)
    .fillna('Other')
    .astype('category')
)

df['detailed household and family stat'].value_counts(dropna=False)

```

Out[16]: detailed household and family stat

Householder	75461
Child	65614
Spouse/Partner	41747
Non-relative/Secondary	6807
Other Relative	6326
Grandchild	3372
Group Quarters	196

Name: count, dtype: int64

```

In [17]: # country
cols = ['country of birth father', 'country of birth mother', 'country of bi
region_map = {
    # --- North America ---
    'United-States': 'North America',
    'Canada': 'North America',

    # --- Latin America / Caribbean ---
    'Mexico': 'Latin America',
    'Puerto-Rico': 'Latin America',
    'Cuba': 'Latin America',
    'Dominican-Republic': 'Latin America',
    'Jamaica': 'Latin America',
    'Honduras': 'Latin America',
    'El-Salvador': 'Latin America',
    'Guatemala': 'Latin America',
    'Colombia': 'Latin America',
    'Ecuador': 'Latin America',
    'Peru': 'Latin America',
    'Nicaragua': 'Latin America',
    'Trinidad&Tobago': 'Latin America',
    'Haiti': 'Latin America',

    # --- Europe ---
    'England': 'Europe',
    'France': 'Europe',
    'Germany': 'Europe',
    'Italy': 'Europe',
    'Poland': 'Europe',
    'Portugal': 'Europe',
    'Ireland': 'Europe',
    'Scotland': 'Europe',
    'Greece': 'Europe',

```

```
'Yugoslavia': 'Europe',
'Hungary': 'Europe',
'Holand-Netherlands': 'Europe',

# --- Asia ---
'China': 'Asia',
'India': 'Asia',
'Japan': 'Asia',
'Philippines': 'Asia',
'Vietnam': 'Asia',
'Korea': 'Asia',
'Cambodia': 'Asia',
'Laos': 'Asia',
'Thailand': 'Asia',
'Taiwan': 'Asia',
'Hong-Kong': 'Asia',

# --- Middle East / Other ---
'Iran': 'Middle East',
'Israel': 'Middle East',
}

for col in cols:
    df[col] = df[col].map(region_map).fillna('Other')
    df[col] = df[col].astype('category')

for col in cols:
    print(f"\n{col} value counts:")
    print(df[col].value_counts())
```

```
country of birth father value counts:
country of birth father
North America    160543
Latin America    18680
Other            8147
Europe          7825
Asia            4095
Middle East      233
Name: count, dtype: int64
```

```
country of birth mother value counts:
country of birth mother
North America    161930
Latin America    18171
Other            7636
Europe          7417
Asia            4171
Middle East      198
Name: count, dtype: int64
```

```
country of birth self value counts:
country of birth self
North America    177689
Latin America    11229
Other            4545
Asia            2975
Europe          2928
Middle East      157
Name: count, dtype: int64
```

```
In [18]: df['class of worker'] = df['class of worker'].replace({
    'Self-employed-not incorporated': 'Self-employed',
    'Self-employed-incorporated': 'Self-employed',
    'Local government': 'Government',
    'State government': 'Government',
    'Federal government': 'Government',
    'Without pay': 'Unemployed',
    'Never worked': 'Unemployed',
    'Not in universe': 'Not in labor force'
}).astype('category')
```

## Feature Engineering

```
In [19]: # Create combined capital income proxy
if all(c in df.columns for c in ['capital gains', 'capital losses', 'dividends from stocks']):
    df['net_capital_income'] = (
        df['capital gains'].fillna(0)
        - df['capital losses'].fillna(0)
        + df['dividends from stocks'].fillna(0)
    )
    df.drop(['capital gains', 'capital losses', 'dividends from stocks'], axis=1)

# Create annual income proxy
if all(c in df.columns for c in ['wage per hour', 'weeks worked in year']):
    df['annual_income_proxy'] = (
```

```

df['wage per hour'].fillna(0)
* df['weeks worked in year'].fillna(0)
* 40 # 40-hour week assumption
)
df.drop(['wage per hour', 'weeks worked in year'], axis=1, inplace=True)

```

```

In [20]: # Domain defined age binning
bins = [0, 17, 25, 35, 50, 65, 100]
labels = ['Child (<18)', 'Young Adult (18-25)', 'Early Career (26-35)',
          'Mid Career (36-50)', 'Late Career (51-65)', 'Senior (65+)']

df['age_group'] = pd.cut(df['age'], bins=bins, labels=labels, right=True, ir
df.drop(columns='age', inplace=True)

```

## Correlation Analysis using Cramer's V

```

In [21]: cat_cols = [c for c in df.columns if df[c].dtype.name == 'category']

```

```

In [22]: def cramers_v(x, y):
          """Compute Cramér's V for two categorical pandas Series."""
          confusion = pd.crosstab(x, y)
          chi2 = chi2_contingency(confusion)[0]
          n = confusion.sum().sum()
          phi2 = chi2 / n
          r, k = confusion.shape
          phi2corr = max(0, phi2 - ((k - 1)*(r - 1)) / (n - 1))
          rcorr = r - ((r - 1)**2) / (n - 1)
          kcorr = k - ((k - 1)**2) / (n - 1)
          return np.sqrt(phi2corr / min((kcorr - 1), (rcorr - 1)))

```

```

In [23]: cramers_results = pd.DataFrame(np.zeros((len(cat_cols), len(cat_cols))),
                                         index=cat_cols, columns=cat_cols)

for i, col1 in enumerate(cat_cols):
    for j, col2 in enumerate(cat_cols):
        if i < j: # upper triangle only
            val = cramers_v(df[col1], df[col2])
            cramers_results.loc[col1, col2] = val
            cramers_results.loc[col2, col1] = val
np.fill_diagonal(cramers_results.values, np.nan)


```

```

In [24]: plt.figure(figsize=(14, 12))
sns.heatmap(
    cramers_results,
    cmap="coolwarm",
    vmin=0,
    vmax=1,
    square=True,
    annot=True, # show correlation numbers
    fmt=".2f", # 2 decimal places
    annot_kws={"size": 7}, # smaller font for readability
    cbar_kws={'shrink': 0.8, 'label': 'Cramér's V'}
)
plt.title("Cramér's V Correlation Between Categorical Variables", fontsize=1

```

```
plt.xticks(rotation=90, fontsize=8)
plt.yticks(rotation=0, fontsize=8)
plt.tight_layout()
plt.show()
```



```
In [25]: threshold = 0.8 # adjust (e.g., 0.75-0.9)
high_corr_pairs = []
for c1 in crammers_results.columns:
    for c2 in crammers_results.index:
        if c1 != c2 and crammers_results.loc[c1, c2] >= threshold:
            high_corr_pairs.append((c1, c2, crammers_results.loc[c1, c2]))

high_corr_df = pd.DataFrame(high_corr_pairs, columns=['Var1', 'Var2', 'Cramers
```

```
In [26]: high_corr_df
```

Out [26]:

	Var1	Var2	CramersV
0	region of previous residence	state of previous residence	0.966
1	state of previous residence	region of previous residence	0.966
2	detailed household and family stat	detailed household summary in household	0.871
3	detailed household summary in household	detailed household and family stat	0.871
4	migration code-change in msa	live in this house 1 year ago	1.000
5	migration code-change in msa	migration prev res in sunbelt	0.824
6	migration code-change in reg	migration code-move within reg	0.869
7	migration code-change in reg	live in this house 1 year ago	1.000
8	migration code-change in reg	migration prev res in sunbelt	0.822
9	migration code-move within reg	migration code-change in reg	0.869
10	migration code-move within reg	live in this house 1 year ago	1.000
11	migration code-move within reg	migration prev res in sunbelt	0.837
12	live in this house 1 year ago	migration code-change in msa	1.000
13	live in this house 1 year ago	migration code-change in reg	1.000
14	live in this house 1 year ago	migration code-move within reg	1.000
15	live in this house 1 year ago	migration prev res in sunbelt	0.992
16	migration prev res in sunbelt	migration code-change in msa	0.824
17	migration prev res in sunbelt	migration code-change in reg	0.822
18	migration prev res in sunbelt	migration code-move within reg	0.837
19	migration prev res in sunbelt	live in this house 1 year ago	0.992
20	country of birth father	country of birth mother	0.824
21	country of birth mother	country of birth father	0.824
22	own business or self employed	veterans benefits	1.000
23	veterans benefits	own business or self employed	1.000

```
In [27]: drop_cols = [
    'migration code-change in msa',
    'migration code-change in reg',
    'migration code-move within reg',
    'migration prev res in sunbelt',
    'detailed household summary in household',
    'region of previous residence',
    'own business or self employed',
    "fill inc questionnaire for veteran's admin",
    'country of birth mother'
```

```
]
df = df.drop(columns=drop_cols)
```

In [28]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 199523 entries, 0 to 199522
Data columns (total 25 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   class of worker                          199523 non-null  category
1   education                                199523 non-null  category
2   enroll in edu inst last wk              199523 non-null  category
3   marital stat                             199523 non-null  category
4   major industry code                     199523 non-null  category
5   major occupation code                   199523 non-null  category
6   race                                    199523 non-null  category
7   hispanic origin                         198649 non-null  category
8   sex                                     199523 non-null  category
9   member of a labor union                 199523 non-null  category
10  reason for unemployment                 199523 non-null  category
11  full or part time employment stat       199523 non-null  category
12  tax filer stat                          199523 non-null  category
13  state of previous residence              199523 non-null  category
14  detailed household and family stat      199523 non-null  category
15  live in this house 1 year ago           199523 non-null  category
16  num persons worked for employer         199523 non-null  category
17  family members under 18                 199523 non-null  category
18  country of birth father                 199523 non-null  category
19  country of birth self                   199523 non-null  category
20  citizenship                             199523 non-null  category
21  veterans benefits                      199523 non-null  category
22  net_capital_income                      199523 non-null  int64
23  annual_income_proxy                     199523 non-null  int64
24  age_group                               199523 non-null  category
dtypes: category(23), int64(2)
memory usage: 7.4 MB
```

## Encoding and Scaling

```
In [29]: edu_order = [['Children',
                      'Below High School',
                      'High School Graduate',
                      'Some College',
                      'Associate Degree',
                      'Bachelors Degree',
                      'Graduate/Professional Degree']]

num_expected = ['annual_income_proxy', 'net_capital_income']
num_cols = [c for c in num_expected if c in df.columns]

use_edu = 'education' in df.columns
```



```
cat_cols = [c for c in df.select_dtypes(include=['category']).columns
             if c != 'education'] # exclude education
```

```
In [30]: transformers = []
if num_cols:
    transformers.append(('num', RobustScaler(), num_cols))
if use_edu:
    transformers.append(('edu', OrdinalEncoder(categories=edu_order), ['education']))
if cat_cols:
    transformers.append(('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat_cols))

preprocessor = ColumnTransformer(transformers)

X = preprocessor.fit_transform(df)
```

```
In [31]: def get_feature_names(ct: ColumnTransformer) -> list:
names = []
for name, transformer, cols in ct.transformers_:
    if hasattr(transformer, 'get_feature_names_out'):
        # e.g., OneHotEncoder
        fn = transformer.get_feature_names_out(cols).tolist()
        names.extend(fn)
    elif name == 'edu' and isinstance(cols, list) and len(cols) == 1:
        names.append(cols[0] + "_ord")
    else:
        # scalars (Robust/Standard) just keep original names
        names.extend(list(cols if isinstance(cols, (list, tuple, np.ndarray)) else [cols]))
return names

feature_names = get_feature_names(preprocessor)
X_df = pd.DataFrame(X, columns=feature_names, index=df.index)

print(f"Encoded+scaled X shape: {X.shape}")
print(f" - numeric scaled: {len(num_cols)}")
print(f" - education ordinal: {int(use_edu)}")
print(f" - categorical OHE block: {X.shape[1] - len(num_cols) - int(use_edu)}")
```

```
Encoded+scaled X shape: (199523, 196)
 - numeric scaled: 2
 - education ordinal: 1
 - categorical OHE block: 193 features
```

## K-selection

```
In [32]: # Sampling data for faster K-selection

rng = np.random.default_rng(42)
n = X.shape[0]
sample_size = min(20000, n)
X_sample = X[rng.choice(n, size=sample_size, replace=False)]
X_sample.shape
```

```
Out[32]: (20000, 196)
```

```

In [33]: # Mini Batch KMeans
Ks = range(3, 11)
Ks = range(3, 11)
results = []

# Run multiple seeds per K to average metrics
for k in Ks:
    inertias, sils, chs = [], [], []
    for seed in [0, 1, 2]:
        km = MiniBatchKMeans(
            n_clusters=k,
            batch_size=4096,
            n_init=20,
            max_iter=300,
            random_state=seed
        ).fit(X_sample)
        labels = km.labels_
        inertias.append(km.inertia_)
        # silhouette on 3K subsample for speed
        idx = rng.choice(sample_size, size=min(3000, sample_size), replace=False)
        sils.append(silhouette_score(X_sample[idx], labels[idx]))
        chs.append(calinski_harabasz_score(X_sample, labels))
    results.append({
        "K": k,
        "Inertia_mean": np.mean(inertias),
        "Inertia_std": np.std(inertias),
        "Silhouette_mean": np.mean(sils),
        "CH_mean": np.mean(chs)
    })

results_df = pd.DataFrame(results)
display(results_df)

```

	K	Inertia_mean	Inertia_std	Silhouette_mean	CH_mean
0	3	1423961888243161.000	315431685492549.438	0.960	35654.934
1	4	1179134453601312.000	296258445085262.312	0.951	29488.214
2	5	1030453749839146.625	37073148578560.703	0.705	24366.879
3	6	1027777600448052.375	3692281960946.654	0.837	19566.200
4	7	971905977081710.500	86568632870658.297	0.866	17385.854
5	8	837049569234161.625	24950665244011.316	0.853	17343.324
6	9	910819833284662.875	70360232060282.031	0.323	13976.788
7	10	995231460840871.375	118132200912435.734	0.426	11800.866

```

In [34]: fig, ax1 = plt.subplots(figsize=(7,4))
ax2 = ax1.twinx()

ax1.errorbar(results_df["K"], results_df["Inertia_mean"],
             yerr=results_df["Inertia_std"], fmt='o-', color='steelblue', capsize=5)
ax2.plot(results_df["K"], results_df["Silhouette_mean"], 's--', color='darkred')


```

```
ax2.plot(results_df["K"], results_df["CH_mean"]/results_df["CH_mean"].max(),

ax1.set_xlabel("Number of Clusters (K)")
ax1.set_ylabel("Inertia (lower better)", color='steelblue')
ax2.set_ylabel("Silhouette / CH (higher better)", color='darkorange')
ax1.set_title("Extended K-Selection – MiniBatchKMeans")
ax1.grid(alpha=0.3)

lines, labels = ax1.get_legend_handles_labels()
lines2, labels2 = ax2.get_legend_handles_labels()
ax1.legend(lines + lines2, labels + labels2, loc='center right')

plt.show()
```



```
In [35]: from kneed import KneeLocator
```

```
kneedle = KneeLocator(
    results_df["K"],
    results_df["Inertia_mean"],
    curve="convex",
    direction="decreasing"
)
auto_k = kneedle.knee
print(f" K = {auto_k}")
```

K = 5

- The elbow in Inertia happens around K = 5 — confirmed by the Kneedle algorithm.
- Silhouette and CH both remain strong up to K ≈ 6–7 but fall sharply beyond that.
- From 3 → 4 → 5, the Silhouette drops because the model starts separating previously cohesive groups. After K = 6, improvements spike and then drop.

**Best estimate: K = 5 or 6**

### Logic for not considering K=8:

1. The biggest drop in Inertia happens between K=3 → K=5, and after that, the curve flattens out. That's a strong indicator that adding clusters beyond 5 gives diminishing returns.
2. Prioritized the CH score over the Silhouette score while choosing since CH tends to remain more stable and robust to local variations, whereas Silhouette can fluctuate due to smaller or uneven clusters.

## Mini Batch K Means

```
In [36]: final_kmeans = MiniBatchKMeans(
    n_clusters=5,
    batch_size=4096,
    n_init=30,
    random_state=42
)
df['cluster'] = final_kmeans.fit_predict(X)
df['cluster'] = df['cluster'].astype('category')
```

## Building Cluster Profiles/ Personas

```
In [37]: # Cluster Sizes
cluster_summary = df['cluster'].value_counts().sort_index()
print("Cluster Sizes:")
print(cluster_summary, "\n")
print(f"Total samples: {len(df)}")
print(f"Cluster Proportions (%):\n{(cluster_summary / len(df) * 100).round(2)}
```

```
Cluster Sizes:
cluster
0      1412
1      2680
2     183927
3       6926
4       4578
Name: count, dtype: int64
```

```
Total samples: 199523
Cluster Proportions (%):
cluster
0      0.710
1      1.340
2     92.180
3      3.470
4      2.290
Name: count, dtype: float64
```

```
In [38]: # Numeric Summaries
num_cols = ['annual_income_proxy', 'net_capital_income']
if all(col in df.columns for col in num_cols):
```

```
num_summary = (
    df.groupby('cluster')[num_cols]
      .agg(['mean', 'median', 'std'])
      .round(2)
)
display(num_summary)
```

	annual_income_proxy			net_capital_income		
	mean	median	std	mean	median	std
cluster						
0	411066.010	418000.000	194929.430	105.940	0.000	1126.550
1	3620695.690	3236480.000	1642130.630	785.280	0.000	4766.730
2	0.180	0.000	44.360	94.990	0.000	674.120
3	1508523.140	1456000.000	443901.770	235.310	0.000	2351.720
4	0.000	0.000	0.000	21264.260	10898.000	27659.970

- Cluster 0: Moderate-income population, likely working-class or lower-middle-income earners with limited investment exposure.
- Cluster 1: Very high earners → average incomes in the millions, moderate capital income. Likely wealthy professionals or high-executive households.
- Cluster 2: Essentially zero income, both wages and capital income are negligible, likely non-working population (children, students, retirees, or not in labor force).
- Cluster 3: Mid-to-upper-income earners with modest capital income, steady full-time workers, professionals, etc.
- Cluster 4: No labor income, but significant investment or capital income → Likely retirees or high-wealth individuals living off investments.

```
In [39]: cat_cols = [c for c in df.select_dtypes(include='category').columns if c !=

def top_categories(series, top_n=3):
    freq = series.value_counts(normalize=True).head(top_n)
    return ', '.join([f"{idx} ({p*100:.1f}%)" for idx, p in freq.items()])

cat_summary = pd.DataFrame({
    col: df.groupby('cluster')[col].apply(top_categories)
    for col in cat_cols
})
display(cat_summary.T)
```

cluster	0	1	2	3	
<b>class of worker</b>	Private (89.9%), Government (10.1%), Unemploye...	Private (78.0%), Government (22.0%), Unemploye...	Not in labor force (53.9%), Private (32.9%), G...	Private (87.9%), Government (12.1%), Unemploye...	Private (45.9%) labor force (
<b>education</b>	Below High School (34.3%), High School Graduat...	High School Graduate (36.9%), Some College (20...	Children (25.8%), High School Graduate (23.5%)...	High School Graduate (42.4%), Some College (24...	Bachelors (
<b>enroll in edu inst last wk</b>	Not in universe (68.5%), High school (18.6%), ...	Not in universe (99.6%), College or university...	Not in universe (93.7%), High school (3.5%), C...	Not in universe (91.5%), College or university...	Not in u (99.3%), Cc univ
<b>marital stat</b>	Never married (57.9%), Married-civilian spouse...	Married-civilian spouse present (69.6%), Never...	Never married (45.0%), Married-civilian spouse...	Married-civilian spouse present (51.9%), Never...	Married-civilian present (
<b>major industry code</b>	Retail trade (40.0%), Education (8.6%), Busine...	Manufacturing-durable goods (15.7%), Hospital ...	Not in universe or children (54.1%), Retail tr...	Retail trade (23.1%), Manufacturing-durable go...	Not in uni children (
<b>major occupation code</b>	Other service (32.2%), Sales (15.4%), Adm supp...	Precision production craft & repair (23.7%), P...	Not in universe (54.1%), Adm support including...	Adm support including clerical (21.5%), Other ...	Not in u (25.2%), Ex admin
<b>race</b>	White (85.7%), Black (10.1%), Amer Indian Aleu...	White (87.5%), Black (8.2%), Asian or Pacific ...	White (83.6%), Black (10.4%), Asian or Pacific...	White (83.4%), Black (11.8%), Asian or Pacific...	White (92.2%) (3.5%), Asian or
<b>hispanic origin</b>	All other (90.6%), Mexican-American (3.0%), Me...	All other (94.0%), Mexican-American (2.2%), Me...	All other (86.0%), Mexican-American (4.2%), Me...	All other (91.1%), Mexican (Mexicano) (3.2%), ...	All other (
<b>sex</b>	Female (59.6%), Male (40.4%)	Male (64.9%), Female (35.1%)	Female (52.7%), Male (47.3%)	Female (56.4%), Male (43.6%)	Male (69.9%),
<b>member of a labor union</b>	No (95.6%), Yes (4.4%), Not in	No (63.5%), Yes (36.5%),	Not in universe (95.9%), No	No (86.8%), Yes (13.2%),	Not in u (88.8%), No Yes

cluster	0	1	2	3	
	universe (0.0%)	Not in universe (0.0%)	(3.5%), Yes (0.5%)	Not in universe (0.0%)	
reason for unemployment	Not in universe (100.0%), Job leaver (0.0%), J...	Not in universe (100.0%), Job leaver (0.0%), J...	Not in universe (96.7%), Other job loser (1.1%...	Not in universe (100.0%), Job leaver (0.0%), J...	Not in u (98.8%), O loser
full or part time employment stat	Children or Armed Forces (52.4%), Full-time sc...	Children or Armed Forces (47.9%), Full-time sc...	Children or Armed Forces (63.1%), Full-time sc...	Children or Armed Forces (50.3%), Full-time sc...	Children or Forces (48.7% ti
tax filer stat	Single (46.8%), Joint both under 65 (30.4%), N...	Joint both under 65 (68.5%), Single (23.3%), H...	Nonfiler (40.7%), Joint both under 65 (32.1%),...	Joint both under 65 (49.5%), Single (38.3%), H...	Joint both u (57.6%), (23.8
state of previous residence	Not in universe (88.0%), Oklahoma (1.1%), Cali...	Not in universe (95.0%), California (0.6%), ? ...	Not in universe (92.0%), California (0.9%), Ut...	Not in universe (91.3%), California (0.6%), Ut...	Not in u (95.7%), C (0.4%
detailed household and family stat	Child (43.6%), Householder (24.2%), Spouse/Par...	Householder (67.2%), Spouse/Partner (24.4%), C...	Householder (36.1%), Child (34.6%), Spouse/Par...	Householder (44.4%), Spouse/Partner (29.5%), C...	Householder ( Spouse/Partner
live in this house 1 year ago	Not in universe under 1 year old (47.6%), Yes ...	Not in universe under 1 year old (52.1%), Yes ...	Not in universe under 1 year old (50.8%), Yes ...	Not in universe under 1 year old (49.7%), Yes ...	Not in universe year old (51.5
num persons worked for employer	6 (33.3%), 1 (20.1%), 2 (15.2%)	6 (51.2%), 4 (16.4%), 3 (10.3%)	0 (51.6%), 6 (16.7%), 1 (11.3%)	6 (35.4%), 4 (18.1%), 3 (14.4%)	6 (31.3%), 0 (2
family members under 18	Not in universe (84.0%), Both parents present ...	Not in universe (100.0%), Both parents present...	Not in universe (70.1%), Both parents present ...	Not in universe (98.4%), Both parents present ...	Not in u (99.9%), Both pr
country of birth father	North America (87.5%), Latin America (5.9%), E...	North America (87.5%), Europe (3.8%), Latin Am...	North America (80.0%), Latin America (9.7%), O...	North America (86.2%), Latin America (6.7%), E...	North , (83.6%), (6.5%), Ot

cluster	0	1	2	3	4
country of birth self	North America (91.4%), Latin America (4.4%), O...	North America (91.5%), Latin America (2.8%), O...	North America (88.9%), Latin America (5.8%), O...	North America (89.8%), Latin America (5.2%), O...	North America (89.8%), Latin America (5.2%), O...
citizenship	Native- Born in the United States (91.1%), For...	Native- Born in the United States (91.3%), For...	Native- Born in the United States (88.6%), For...	Native- Born in the United States (89.5%), For...	Native- Born in the United States (89.5%), For...
veterans benefits	2 (99.4%), 1 (0.6%), 0 (0.0%)	2 (99.0%), 1 (1.0%), 0 (0.0%)	2 (73.2%), 0 (25.8%), 1 (1.0%)	2 (99.5%), 1 (0.5%), 0 (0.0%)	2 (98.5%), 1 (1.5%), 0 (0.0%)
age_group	Young Adult (18–25) (37.5%), Early Career (26–35) (62.5%)	Mid Career (36–50) (50.3%), Early Career (26–35) (49.7%)	Child (<18) (30.3%), Mid Career (36–50) (20.5%), Early Career (26–35) (49.2%)	Mid Career (36–50) (30.1%), Early Career (26–35) (69.9%)	Mid Career (36–50) (38.8%), Late Career (51–65) (61.2%)

**Cluster 0:** young adults (18–25), female-majority, employed in private sector (90%), mainly in retail trade or service occupations.

**Cluster 1:** Mostly men (65%), mid-career (36–50), employed in private/government sectors with manufacturing and hospital services dominance

**Cluster 2:** Large share of children (< 18), not in labor force (54%), mostly female, and never married.

**Cluster 3:** Mixed gender, early- to mid-career adults (26–50), mostly private employees in clerical / administrative / service roles.

**Cluster 4:** Predominantly men (70%), late-career to senior (50+), private or self-employed, many householders, few dependents.

## Based on the above numeric and categorical composition analysis, these are the customer segments:

1. Cluster 0 — Young Working-Class Employees
2. Cluster 1 — Mid-Career Married Professionals
3. Cluster 2 — Non-Earning / Dependent Population
4. Cluster 3 — Mid-Income Dual-Worker Households



## 5. Cluster 4 — Educated, High-Wealth / Retired Investors

```
In [41]: persona_data = [
    {
        "Cluster": "0 – Young Working-Class Employees",
        "Economic Tier": "Lower income",
        "Employment Type": "Private wage earners",
        "Profile Summary": (
            "Young adults (18–25), mainly female, employed in retail/service",
            "Low education (Below HS/HS Grad), never married, low income (~4"
        )
    },
    {
        "Cluster": "1 – Married Skilled Professionals",
        "Economic Tier": "Upper-middle income",
        "Employment Type": "Private / Government professionals",
        "Profile Summary": (
            "Mid-career men (36–50), married, high wages (~3.6M), modest cap",
            "Educated, steady jobs in manufacturing and hospital sectors."
        )
    },
    {
        "Cluster": "2 – Dependents / Non-workers",
        "Economic Tier": "Low / inactive",
        "Employment Type": "Not in labor force",
        "Profile Summary": (
            "Children, students, or unemployed adults; minimal or zero incom",
            "Largely dependents supported by others in household."
        )
    },
    {
        "Cluster": "3 – Steady Working Households",
        "Economic Tier": "Middle income",
        "Employment Type": "Private / Clerical / Service roles",
        "Profile Summary": (
            "Early-to-mid-career adults (26–50), mix of genders, high-school",
            "stable full-time employment (~1.5M income), modest investments."
        )
    },
    {
        "Cluster": "4 – Educated Asset-Rich Retirees",
        "Economic Tier": "High wealth",
        "Employment Type": "Self-employed / Retired investors",
        "Profile Summary": (
            "Older men (50+), highly educated (Bachelor's+), married, ",
            "no wage income but high capital income (~21K). Financially inde"
        )
    }
]

# Convert to DataFrame
persona_df = pd.DataFrame(persona_data)

# === Styling for Notebook Display ===
def color_tier(val):
```

```

if "Low" in val: return "background-color:#ffe6e6"
if "Middle" in val: return "background-color:#fff5cc"
if "Upper" in val: return "background-color:#e6f3ff"
if "High" in val: return "background-color:#d6f5d6"
return ""

styled = (
    persona_df.style
        .set_properties(**{
            'text-align': 'left',
            'white-space': 'pre-wrap',
            'font-size': '13px'
        })
        .applymap(color_tier, subset=["Economic Tier"])
        .hide(axis='index')
        .set_table_styles([
            {"selector": "th", "props": [("font-size", "13px"), ("text-align", "left")]},
            {"selector": "td", "props": [("border", "1px solid #ddd"), ("padding", "5px 15px 5px 5px")]}
        ])
)
display(styled)

```

Cluster	Economic Tier	Employment Type	Profile Summary
0 – Young Working-Class Employees	Lower income	Private wage earners	Young adults (18–25), mainly female, employed in retail/service sectors. Low education (Below HS/HS Grad), never married, low income (~400K), minimal assets.
1 – Married Skilled Professionals	Upper-middle income	Private / Government professionals	Mid-career men (36–50), married, high wages (~3.6M), modest capital income. Educated, steady jobs in manufacturing and hospital sectors.
2 – Dependents / Non-workers	Low / inactive	Not in labor force	Children, students, or unemployed adults; minimal or zero income and assets. Largely dependents supported by others in household.
3 – Steady Working Households	Middle income	Private / Clerical / Service roles	Early-to-mid-career adults (26–50), mix of genders, high-school educated, stable full-time employment (~1.5M income), modest investments.
4 – Educated Asset-Rich Retirees	High wealth	Self-employed / Retired investors	Older men (50+), highly educated (Bachelor's+), married, no wage income but high capital income (~21K). Financially independent households.

## Future Scope

1. Although *MiniBatch K-Means* produced clear and interpretable clusters, exploring other clustering algorithms like DBSCAN, Hierarchical Clustering, etc. could help uncover more complex or non-linear group structures.
2. Incorporating PCA (Principal Component Analysis) before clustering can reduce redundancy among correlated features, enhance cluster stability, and allow for clear 2D/3D visualizations of group separation.
3. Future work can also include statistical validation (ANOVA, Chi-Square) and predictive modeling, training a classifier to automatically assign new individuals to socio-demographic segments based on their characteristics.

In [ ]: